

An Efficient Ear Decomposition Algorithm ^{*}

Debarshi Dutta ^{*1}, Kishore Kothapalli ^{†1}, G. Ramakrishna ^{‡§2}, Sai Charan Regunta ^{¶1}, and Sai Harsh Tondomker ^{||1}

¹International Institute of Information Technology, Hyderabad. India.

²Department of Computer Science and Engineering, Indian Institute of Technology, Tirupati. India.

Abstract

An ear decomposition of a graph G is a partition of the edge set of G into a sequence of edge-disjoint paths, such that only the end vertices of each path appear on earlier paths. For a graph on n vertices and m edges, the state-of-art algorithm for obtaining an ear decomposition by Schmidt takes $O(m + n)$ time. We design and implement an algorithm, using Schmidt's algorithm, to obtain an ear decomposition for a biconnected graph, whose running time is $O(m + n)$. In practice, however, our experiments reveal that, the proposed algorithm runs at least 2 times faster than Schmidt's algorithm. The speedup increases as the graph gets denser.

1 Introduction

Obtaining an ear decomposition of a graph is an important problem in the context of graph algorithms. An ear decomposition of a graph is used in several other graph algorithms such as testing connectivity, s - t -numbering and planarity-testing [9]. Also, an ear decomposition has been used as a paradigm to obtain parallel algorithms for various problems. Whitney has introduced the notion of ear decomposition to characterize biconnected graphs [10]. An *ear* of a graph G is a maximal path whose internal vertices have degree 2. An *ear decomposition* of a graph $G = (V, E)$ is a partition of E into a sequence (P_0, P_1, \dots, P_k) , such that (i) P_0 is a cycle, (ii) for each $i \geq 1$, P_i is an ear of $P_0 \cup \dots \cup P_{i-1}$. An ear decomposition (P_0, P_1, \dots, P_k) of G is an *open ear decomposition* if the end points of all the ears P_i , $i \geq 1$, are distinct.

Past Work. The concept of ear decomposition is very well studied both structurally and algorithmically. An ear decomposition is used to characterize biconnected graphs [10]. Further, the notion of nested ear decomposition and nice ear decompositions are developed to characterize series-parallel graphs and polygonal 2-trees, respectively [3, 7]. For a graph on n vertices and m edges, algorithms to obtain an ear decomposition in $O(m + n)$ time are known for a long time. Lovász has designed an algorithm for the first time to obtain an ear decomposition in parallel framework [5]. The state-of-art serial algorithm to construct an ear decomposition is proposed by Schmidt [9]. His algorithm is based on depth first search spanning tree and is simpler to visualize. The algorithm of Schmidt also runs in time $4m + O(n)$.

Motivation and Our Contribution. To the best of our knowledge, there are no studies on computing ear decomposition from a practical perspective in serial computing. The state of art algorithm by Schmidt is simple and elegant [9]. However, this algorithm has not been explored in practice. The main objective

^{*}debarshi.dutta@research.iiit.ac.in

[†]kkishore@iiit.ac.in

[‡]rama@iittp.ac.in

[§]Work initiated when the autor was at the Indian Institute of Information Technology, Sri City. India.

[¶]saicharan.regunta@research.iiit.ac.in

^{||}sai.harsh@research.iiit.ac.in

^{*}A preliminary version of this paper appeared as *An Efficient Ear Decomposition Algorithm* in the proceedings of CTW 2017.

of this work is to obtain an ear decomposition algorithm that works well both in theory and practice. In this paper, we present an $O(m + n)$ algorithm that offers an improvement to the algorithm of Schmidt in the practical setting. In particular, we show that, a large number of edges of a graph are *redundant* in the process of obtaining an ear decomposition. Removing such redundant edges in a prior pre-processing step can often result in practical improvement. Our characterization of redundant edges (trivial ear) is based on a similar notion that is identified in the context of biconnectivity [1]. In practice, our algorithm runs at least 2 times faster than Schmidt's algorithm on graphs having $\Omega(n \log n)$ edges.

2 Algorithm for Ear Decomposition

We use standard graph terminology from [10]. For a graph G , let $n = |V(G)|$ and $m = |E(G)|$ denote the number of vertices and edges in G , respectively. A graph is *biconnected* (*2-vertex connected*) if it does not contain a cut-vertex. An edge e in a biconnected graph is *non-essential* if the graph remains biconnected after the removal of e [1]. For $i \geq 2$, an ear P_i in an ear decomposition is called a *trivial ear* if the number of edges in P_i is one. In the rest of the paper, G denotes an unweighted and undirected biconnected graph.

Schmidt has provided a simple and elegant linear-time algorithm to verify the 2-vertex connectivity and 2-edge connectivity of a graph [9]. We shortly describe Schmidt's algorithm to obtain an ear decomposition of a biconnected graph G . The first step in this algorithm is to obtain a depth first search tree T of G . During the construction of T , a *depth first search* (DFS) number $\text{DFS}[u]$, is maintained at every vertex u , that indicate the visiting time of u . Further, we have $m - n + 1$ number of iterations that produce an ear in each iteration. Let $e(r, v)$ be a non-tree edge incident at root vertex r in T . Then the cycle formed with the non-tree edge (r, v) and the tree path between r and v , is a first ear P_0 . We then mark all the vertices in P_0 as visited. The rest of the $m - n$ non-tree edges (u, v) , such that $\text{DFS}[u] < \text{DFS}[v]$ are processed as follows in non-decreasing order based on $\text{DFS}[u]$. For each non-tree edge $e = (u, v)$, we traverse T from v in upward direction till a visited vertex is encountered; then the edge (u, v) followed by the sequence of edges in the upward path forms a new ear. Also, all the vertices that are traversed are marked as visited in each iteration. The time required to construct a depth first search tree takes $2m + O(n)$. For each non-tree edge (u, v) , if $\text{DFS}[u] < \text{DFS}[v]$, then the non-tree edge is processed; otherwise, (u, v) is ignored. However a comparison is required in both cases. Thus the total time to obtain an ear decomposition is $4m + O(n)$.

Algorithm 1 shows an overview of our approach to compute an ear decomposition. At a higher-level, the main idea in our approach is to filter many non-essential edges and compute an ear decomposition on the rest of the graph.

Algorithm 1: An algorithm to find an ear decomposition of a biconnected graph G

- 1 Construct a breadth first search (BFS) spanning tree T of G ;
 - 2 Construct a spanning forest F from G' , where $G' = G - T$;
 - 3 Find an ear-decomposition \mathcal{P} of $T \cup F$ using Schmidt's algorithm ;
 - 4 return the sequence of ears in \mathcal{P} and the edges in G'' as trivial ears, where $G'' = G' - F$;
-

First we construct a spanning tree T of G by applying a breadth first search algorithm. Next we obtain a spanning forest F from $G - T$. We then compute an ear-decomposition \mathcal{P} of $T \cup F$ using Schmidt's algorithm. Further, append all the edges in $G - (T \cup F)$ to \mathcal{P} as trivial ears to obtain an ear decomposition of the original graph. The correctness of this algorithm follows from Theorem 2.3.

We use the following lemma on the characterization of biconnected graphs to prove Theorem 2.3.

Lemma 2.1 ([10]) *A graph G is biconnected (2-vertex connected) if and only if G has an open ear decomposition.*

The following lemma helps to find the number of biconnected components in a connected graph, and is useful in the proof of Theorem 2.3.

Lemma 2.2 ([1]) *Let T be a BFS spanning tree of G and F be a spanning forest in $G - T$. Then, the edges of each connected component of $G - T$ are in one biconnected component. The number of biconnected components in G and $T \cup F$ is same.*

Theorem 2.3 *For a biconnected graph G , let T be a BFS-spanning tree of G and F be a spanning forest in $G - T$. Then there is an ear decomposition \mathcal{P} of G in which every edge in $G - (T \cup F)$ is a trivial ear.*

Proof From Lemma 2.2, $T \cup F$ is biconnected. Then, by Lemma 2.1, there is an open ear decomposition \mathcal{P}' of $T \cup F$. As a result, \mathcal{P}' with each edge in $G - (T \cup F)$ being a trivial ear, becomes an open ear decomposition of G with the mentioned property. \square

In Algorithm 1, Line 1 and Line 3 take $2m + O(n)$ time and $O(n)$ time, respectively. Constructing a spanning forest F from $G - T$ efficiently in Line 2 is a more involved task. If we use breadth first or depth first traversals, then Line 2 consumes $2m + O(n)$ time, and hence, the total running time is $4m + O(n)$. This matches with Schmidt's algorithm, and hence this is not a promising idea. If we use disjoint-set-forest data structure with union-by-rank and path-compression, Line 2 consumes $m\alpha(n) + O(n)$ time, where $\alpha(n)$ is a very slowly growing function. However, $\alpha(n) = 3$ for $8 \leq n \leq 2048$ [2]. This means that the total run time exceeds that of Schmidt's algorithm even for fairly small values of n . A linked-list based disjoint-set data-structure requires $m + O(n \log n)$ time to compute the task in Line 2. However, as data-locality is very poor in linked-lists, our experiments reveal that this idea does not beat Schmidt's algorithm in practice.

To make our idea work, we introduce a randomized algorithm Algorithm 2, shown below for the implementation of Line 2 in Algorithm 1.

Algorithm 2: Algorithm to implement Line 2 in Algorithm 1

```

1  $X = \{v \in V(G) \mid \text{degree}(v) \text{ in } T = \text{degree}(v) \text{ in } G\}, F = \emptyset$  ;
2 for each vertex  $u \in V(G) - X$  do
3    $\lfloor$  append each edge incident at  $u$  to  $F$  with probability  $\frac{\log n}{n}$  ;

```

The key idea in the randomized algorithm is to choose a sparse spanning connected subgraph of $G - X$, instead of a spanning forest of G . For each vertex $u \in V(G)$, each edge incident at u is sampled with probability $\frac{\log n}{n}$ and all the sampled edges in the random process form a spanning connected subgraph of $G - X$. This insight is inferred from our experiments and support the correctness of Algorithm 2. Let the number of vertices and edges in the input graph H be n vertices and $\Omega(n \log n)$ edges, respectively. If H' is a spanning subgraph of H that is constructed by sampling each edge of H with probability $\frac{\log n}{n}$, then we claim that H' is connected with high probability. The proof of this claim is shown for random graphs, in Lemma 2.5 with the help of Lemma 2.4. The expected number of edges in F is clearly in $O(n \log n)$. If we also ensure in the sampling that each edge uv is sampled exactly once, then the number of edges in F is also in $O(n \log n)$ with high probability as can be shown using Chernoff bounds [6]. If we assume that each edge can be sampled in constant time, then Algorithm 2 runs in $m + O(n \log n)$ time, and hence the run time of Algorithm 1 is $3m + O(n \log n)$. This run time is further reduced to $2m + O(n \log n)$ time in next section. Since m is $\Omega(n \log n)$, our algorithm runs in $O(m + n)$ time. However, our algorithm performs well in practice as m increases beyond $\Omega(n \log n)$.

The above paragraph establishes that F is sparse. To show that F is also connected, we use the following ideas from random graphs.

Lemma 2.4 [4] *A random graph $G(n, p)$ on n vertices and edge probability p is connected with very high probability, if $p = \frac{\ln n}{n}$.*

Since a necessary condition for a graph to have an ear decomposition is for the graph to be biconnected, in the following, we consider graphs that are biconnected. A random graph $G(n, p)$ is biconnected provided $p \geq \frac{1}{n} \cdot \left(\log n + \frac{\beta}{2} \cdot \log \log n + \alpha + o(n) \right)$, for constants α and β as shown by Reif and Spirakis in [8].

Lemma 2.5 *Let G be a biconnected random graph $G(n, p)$ on n vertices, in which each edge is included with*

probability p . Let G' be the graph obtained from G , by sampling each edge in G with probability $\frac{\ln n}{pn}$. Then G' is connected with high probability.

Proof Let H be an arbitrary graph on n vertices and let $\{e_{i1}, \dots, e_{ik}\}$ be the set of edges in H . For each edge $e \in E(H)$, the probability of e available in G' is $p \times \frac{\log n}{p \cdot n} = \frac{\ln n}{n}$. For each edge $e \notin E(H)$, the probability of e not available in G' is $(1 - p) + p \times (1 - \frac{\ln n}{p \cdot n}) = 1 - \frac{\ln n}{n}$. From these two observations, we can derive that the probability of $E(H)$ equals to $E(G')$ is $P(E(H) = E(G')) = (\frac{\ln n}{n})^k \times (1 - \frac{\ln n}{n})^{\binom{n}{2}-k}$. Let G'' be a random graph chosen on n vertices with edge probability $\frac{\ln n}{n}$. The probability of $E(H)$ equals to $E(G'')$ is same as $P(E(H) = E(G))$. From Lemma 2.4, G'' is connected with high probability, and hence G' is connected with high probability. \square

3 Implementation Details and Experiments

In this section, we describe the lower level details of our algorithm. Without loss of generality, let us assume that the vertices in the input graph are numbered from 1 to n . We use the *compressed sparse row* (CSR) representation to store the input graph and as well as the intermediate graphs. A CSR of a graph G consists of two arrays, namely a vertex array $V[\]$ and an edge array $E[\]$. For each i , $1 \leq i \leq n$, $V[i]$ denotes an index in $E[\]$ such that all the neighbors of i are stored in $E[\]$ at locations $V[i], \dots, V[i+1] - 1$.

We recall the construction of a spanning forest F in Algorithm 2. Each edge incident at a vertex u is sampled with probability $\frac{\log n}{n}$. Then, the expected number of edges in F that are incident on u for a dense graph is at most $\log n$. In practice, this step consumes more time as every edge needs to be processed. One way to perform this step is as follows. For sampling k edges that are incident on a vertex i , we first choose k random indices whose range is $V[i], \dots, V[i+1] - 1$, and move these edges from G to F . If the number of edges incident at a vertex are less than k , then we move all of them from G to F . The run-time of Algorithm 2 as per this idea is $O(n \log n)$, and thus the run-time of our main algorithm is $2m + O(n \log n)$. The edges that are not sampled from CSR of G will be remained as trivial ears.

Experimental Results

We implement our algorithms in C using gcc version 4.8.5 compiler with -O3 optimization level. We use Intel(R) Xeon(R) CPU E5-2650 v3 processor, which is based on x86_64 architecture. The frequency of the processor that we use in our experiments is 2.30GHz. The details on the number of vertices and density (m/n) of the graphs considered in our experiments, which are constructed using GT-generator, are given in the figures shown below.




Figure 1: Density Vs SpeedUp

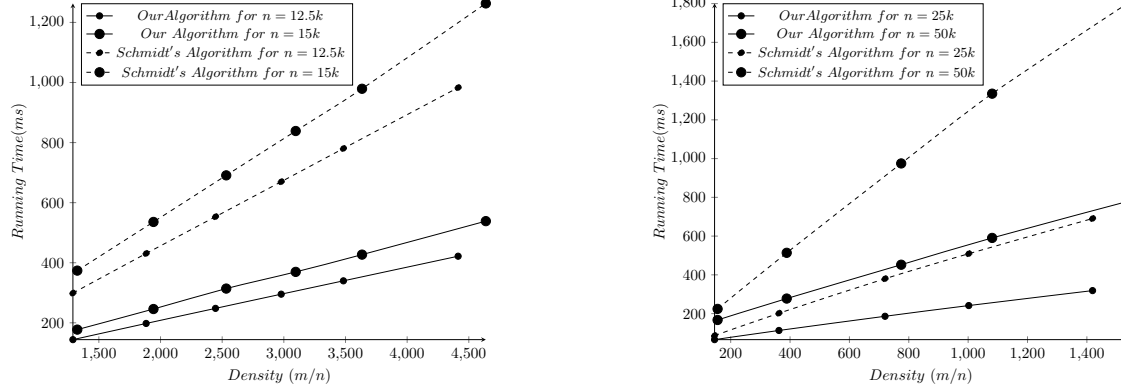


Figure 2: Schmidt's Vs Our Algorithm

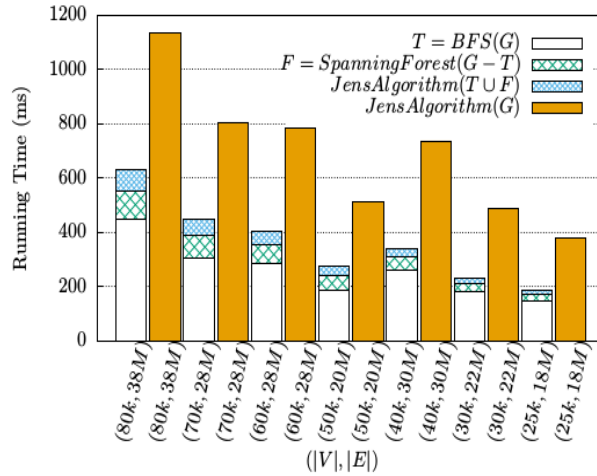


Figure 3: Profiling

Our experiments reveal that the speedup is proportionate to the density of a graph. In other words, as the density of a graph increases, the speedup of our algorithm increases further and this is shown in the left-side figures. The run-time of our algorithm against Schmidt's algorithm is shown in Figure 2. In practice, our algorithm runs at least 2 times faster than Schmidt's algorithm. The execution times of Line 1, Line 2 and Line 3 in Algorithm 1, along with the execution time of Schmidt's Algorithm, on graphs of various sizes, are shown in Figure 3. Since many edges are filtered in our algorithm, the execution times of Line 2 and Line 3 are significantly smaller than that of Line 1.

Remark. Our algorithm can be used to obtain an approximated minimal biconnected graph. Finding the trade off between the quality of the solution and runtime by our algorithm against the state-of-art approximation algorithms to find a minimal biconnected graph is an interesting study.

References

- [1] G. Cong and D. A. Bader. An experimental study of parallel biconnected components algorithms on symmetric multiprocessors (smps). *Inter. Par. and Dist. Proc. Symp.*, 2005.
- [2] Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, and Clifford Stein. *Introduction to Algorithms, Third Edition*. The MIT Press, 3rd edition, 2009.

- [3] D. Eppstein. Parallel recognition of series-parallel graphs. *Inf. Comput.*, 98(1):41–55, 1992.
- [4] P. Erdős and A Rényi. On the evolution of random graphs. *Publ. Math. Inst. Hungar. Acad. Sci*, pages 17–61, 1960.
- [5] L. Lovász. Computing ears and branchings in parallel. *Found. of Comp. Sci*, 464–467, 1985.
- [6] R. Motwani and P. Raghavan. *Randomized Algorithms*. Cambridge University Press, 1995.
- [7] N.S. Narayanaswamy and G. Ramakrishna. On minimum average stretch spanning trees in polygonal 2-trees. *Theor. Comput. Sci.*, 575(C):56–70, 2015.
- [8] J. H. Reif, , and P. G. Spirakis. k-connectivity in random undirected graphs. *Discrete Mathematics*, pages 181 –191, 1985.
- [9] J. M. Schmidt. A simple test on 2-vertex- and 2-edge-connectivity. *Info. Proc. Lett.*, 113(7):241–244, 2013.
- [10] Douglas B. West. *Introduction to graph theory - second edition*. Prentice Hall, 2001.